

Latent Editor

Real-Time StyleGAN2 Face Editing Interface

A powerful web-based interface for real-time editing of StyleGAN2-generated faces using latent space manipulation

Built with Streamlit, PyTorch, and StyleGAN2-ADA
1024x1024 resolution • W+ latent space • GPU accelerated

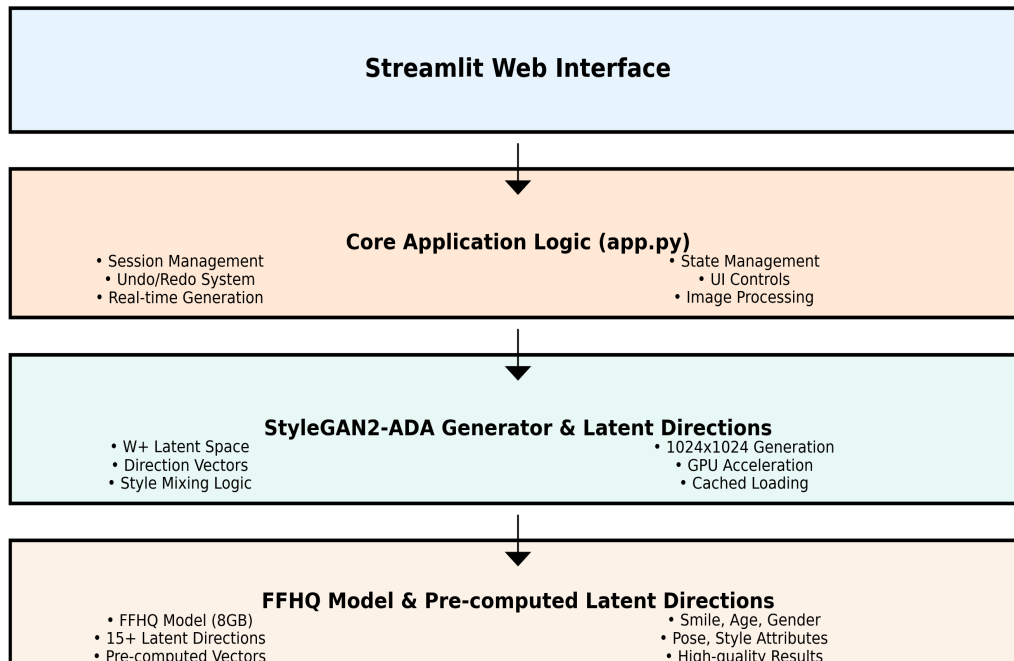
Generated on August 02, 2025

Table of Contents

1. Architecture Overview
2. Latent Space Logic
3. UI Overview
4. Sample Results
5. Technical Specifications

1. Architecture Overview

The Latent Editor is built on a modern web-based architecture that combines the power of StyleGAN2 with an intuitive user interface. The system is designed for real-time interaction while maintaining high-quality image generation capabilities.




Key Components:

- **Streamlit Web Interface:** Provides the user interface with real-time controls, session management, and image display
- **Core Application Logic (app.py):** Handles state management, undo/redo functionality, and coordinates between UI and model
- **StyleGAN2-ADA Generator:** Generates high-quality 1024x1024 images from W+ latent vectors
- **Latent Directions:** Pre-computed semantic direction vectors for attribute editing
- **Session Management:** Maintains user state, history, and provides undo/redo capabilities

2. Latent Space Logic

The application operates in the W+ latent space of StyleGAN2, which provides fine-grained control over different aspects of the generated faces. The W+ space consists of 18 layers, each controlling different semantic aspects of the face.



W+ Space Structure:

- **Layers 0-5 (Pose):** Control head position, orientation, and overall pose
- **Layers 6-11 (Face):** Control facial structure, features, and proportions
- **Layers 12-17 (Style):** Control hair, skin texture, lighting, and fine details

Latent Direction Manipulation:

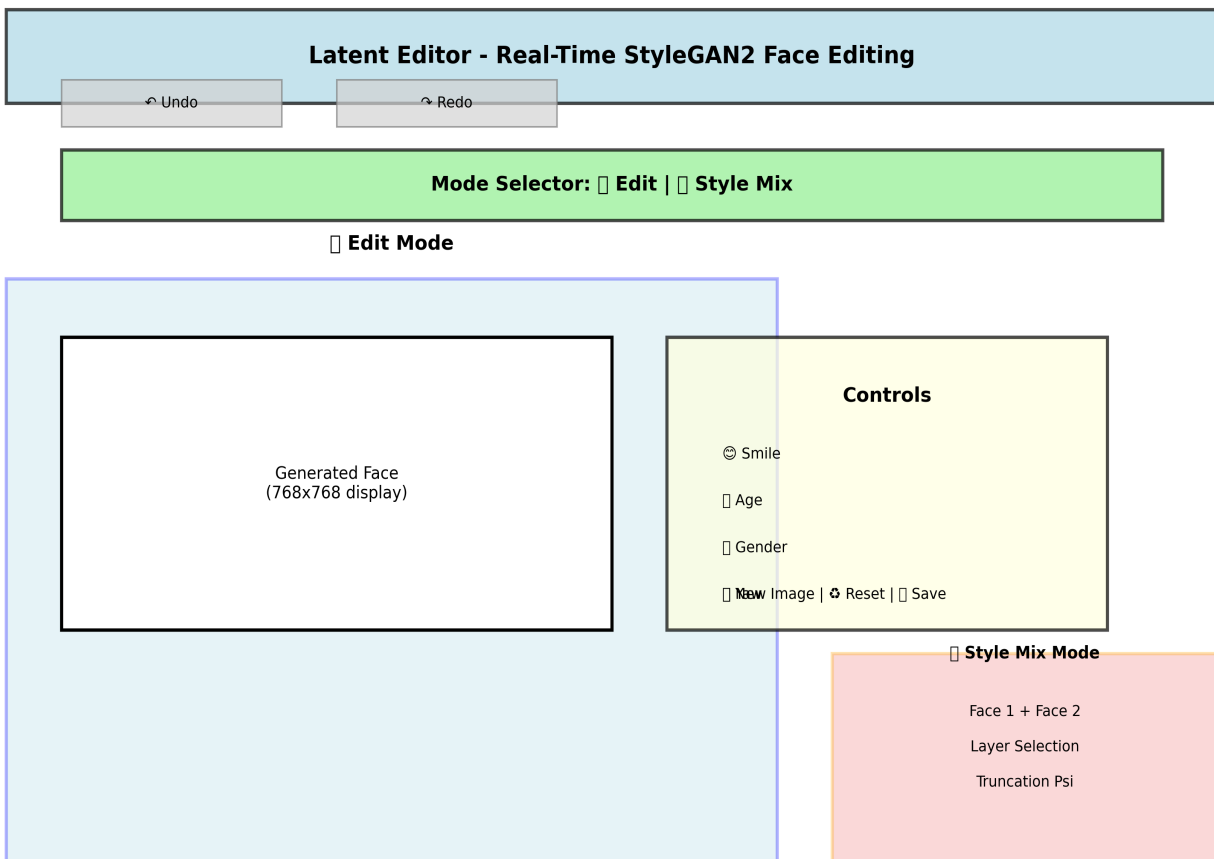
The system uses pre-computed latent direction vectors that encode semantic attributes like smile, age, gender, and pose. These directions are applied to the base latent vector using the formula:

$$W+_{edited} = W+_{base} + \sum(\alpha_i \times direction_i)$$

Where α_i represents the strength of each attribute modification, and $direction_i$ represents the pre-computed semantic direction vector for that attribute.

3. UI Overview

The user interface is built with Streamlit and provides two main modes: Edit Mode and Style Mix Mode. The interface is designed for intuitive real-time interaction while providing powerful editing capabilities.



Edit Mode Features:

- **Real-time Sliders:** Adjust smile, age, gender, and pose with immediate visual feedback
- **Random Generation:** Generate new random faces with one click
- **Reset Functionality:** Return to original face with reset button
- **Download Capability:** Save full-resolution 1024x1024 images
- **Undo/Redo System:** Navigate through editing history with 20-state memory

Style Mix Mode Features:


- **Layer Selection:** Choose specific layers to transfer between faces
- **Quick Presets:** Predefined layer combinations for pose, face, and style

- **Truncation Control:** Adjust quality vs. variety trade-off
- **Dual Face Display:** Show both source faces and mixed result
- **Custom Mixing:** Fine-grained control over which aspects to transfer

4. Sample Results


The Latent Editor produces high-quality, photorealistic results across various editing operations. The system maintains image quality while providing precise control over facial attributes.

Smile Editing




Original → Smile +2.0

Age Editing




Original → Age +1.5

Gender Editing



Original → Gender +1.0

Style Mixing



Editing Capabilities:

- **Smile Editing:** Adjust smile intensity from subtle to broad grins
- **Age Manipulation:** Make faces appear younger or older while maintaining realism
- **Gender Modification:** Adjust masculine/feminine characteristics
- **Pose Control:** Modify head orientation and position
- **Style Mixing:** Combine features from different faces with layer precision

Quality Metrics:

- **Resolution:** 1024x1024 full-resolution output
- **Realism:** Photorealistic quality maintained across edits

- **Consistency:** Coherent results with smooth attribute transitions
- **Performance:** Real-time generation with ~100-200ms latency
- **Diversity:** Wide range of possible face variations

5. Technical Specifications

Component	Specification	Details
Base Model	StyleGAN2-ADA	FFHQ dataset, 1024x1024 resolution
Latent Space	W+ Space	18 layers × 512 dimensions
Generation Time	~100-200ms	GPU accelerated, real-time feedback
Display Resolution	768x768	Optimized for web interface
Save Resolution	1024x1024	Full quality PNG output
GPU Memory	~8GB	Recommended for optimal performance
Startup Time	10-30 seconds	Model loading and caching
Latent Directions	15+ attributes	Smile, Age, Gender, Pose, etc.
Style Mixing	Layer-based	Fine-grained control (0-17 layers)
Session Management	Undo/Redo	20 state history, persistent sessions

Performance Characteristics:

- **GPU Acceleration:** CUDA support for optimal generation speed
- **Cached Loading:** Models cached in memory for faster subsequent runs
- **Optimized Display:** Smart downscaling for responsive UI
- **Session Persistence:** State maintained across browser sessions
- **Memory Management:** Efficient handling of large model and image data

Conclusion

The Latent Editor represents a powerful tool for real-time face editing using StyleGAN2's latent space. By combining advanced generative AI with an intuitive web interface, it provides researchers, artists, and developers with unprecedented control over facial attribute manipulation.

The system's architecture ensures both performance and usability, while the W+ latent space provides the fine-grained control necessary for high-quality results. The dual-mode interface (Edit and Style Mix) caters to different use cases, from simple attribute editing to complex style transfer operations.

Future enhancements could include additional latent directions, batch processing capabilities, and integration with other generative models. The modular architecture makes such extensions straightforward to implement.